

SecondBrain – v0.1 MVP Spec (4-week Test Sprint)

Version: Draft v0.1

Founders:

- Nour – product/clinical
- Saahil – tech/implementation

Timeframe: 4 weeks (part-time for both)

1. Goal of v0.1

Build a **very small, usable first version** that:

- A resident / busy physician can actually use during real shifts
- Helps them **organize patients and tasks into a single mental queue**
- Gives us **real usage data and feedback** from at least 3–5 doctors

If, at the end of 4 weeks, 3–5 doctors say:

“This is actually useful, I’d like to keep using it if you improve it,”
then v0.1 is a success.

2. Constraints (Reality Check)

- Both of us are **part-time** for the next 1–2 months (IELTS, current project, etc.)
 - No hospital integrations in v0.1
 - No heavy AI/LLM features yet
 - Must be **simple web app**, mobile-friendly (most doctors will use it on phone)
 - Focus on **one core workflow**, not many features
-

3. Target User Persona for v0.1

“Internal medicine resident in a busy hospital”

- 10–30 inpatients under their care
- Constantly adding mental tasks:
 - “Check K+ in 4 hours”
 - “Follow up CT report”
 - “Call cardiology back”
 - “Review antibiotics in the evening”
- Juggles:
 - Rounds

- New admissions
- Calls / messages
- Documentation
- Biggest feeling:
“I’m forgetting something important...”

We design v0.1 for this person, not for everyone.

4. Core Use Case for v0.1

Scenario:

- Resident starts / is in the middle of their shift
- Wants to see, in one place:
 - All their patients
 - All open tasks
 - What they should focus on in the next 20–60 minutes

v0.1 must answer one question well:

“Given my patients and tasks, what should I work on now?”

No AI magic is required yet – smart sorting + a clean UI already beats their current chaos.

5. v0.1 Scope – Features

5.1. Entities

We keep the data model very simple:

- **User**
 - id
 - name
 - email (optional in v0.1)
- **Patient**
 - id
 - user_id (owner doctor)
 - name / nickname
 - bed / room (optional free text)
 - main problem / diagnosis (free text)
 - priority flag: normal / important / critical
 - status: active / discharged
- **Task**
 - id
 - patient_id

- o title (short text: “Recheck K+”, “Call cardio”, ...)
- o notes (optional longer text)
- o due time (datetime or simple “today / tomorrow”)
- o priority: low / medium / high
- o status: open / done
- o created_at, completed_at

No more than this for v0.1.

5.2. Screens / Flows

Screen 1 – Today Queue (Home)

This is the **main screen**.

- Shows **list of open tasks across all active patients**
- Each row:
 - o Patient name (with small badge if critical)
 - o Task title
 - o Due time (or “Overdue / Today / Tomorrow”)
 - o Priority indicator
 - o Checkbox / button to mark as “Done”

Sorting logic v0.1:

1. Overdue tasks first
2. Then by due time (earlier first)
3. Then by priority (high > medium > low)

Basic filters (optional but nice if easy):

- Filter by patient
- Filter by priority
- Toggle: “Only critical patients”

From this screen user can:

- Tap task → go to patient detail
 - Mark task as done
 - Add new task (with patient selection)
-

Screen 2 – Patients List

Shows all **active patients** of that user.

- For each patient:

- Name
- Bed / room
- Main problem
- Priority badge (normal / important / critical)
- Count of open tasks (e.g. “3 open tasks”)

User can:

- Add new patient
 - Tap patient → go to Patient Detail
-

Screen 3 – Patient Detail

For each patient:

- Header:
 - Name
 - Bed / room
 - Main problem
 - Priority flag (user can change)
- Section: **Open Tasks**
 - List of open tasks for this patient (title, due time, priority, mark done)
- Section: **Completed Tasks (collapsed by default)**
- Section: **“What changed” (manual v0.1)**
 - For now, just:
 - A free text area: “Notes since last review”
 - Or a simple list of recent edits (v0.1 can be very basic)

User can:

- Add new task for this patient
 - Edit patient info
 - Change patient priority
 - Mark patient as “discharged” (moves out of active list, but not deleted)
-

Screen 4 – (Very) Simple Onboarding / Access

v0.1 must be **easy to give to test doctors**, so:

Option A (simplest for test):

- No full auth yet
- One shared “test link” per doctor (with a unique URL token)
- Each token = one “User”

Option B (if not hard):

- Simple email + password registration
- Or magic-link login

We can decide based on your tech preference, but key is:

- **Speed to onboard test users**
 - Not security perfection in v0.1
-

5.3. Basic Analytics / Tracking (for us only)

We need minimal tracking to learn:

- How many times a user logged in
- How many patients created
- How many tasks created / completed
- When they last used it

So v0.1 should at least store:

- `last_login_at` for each user
- `created_at` / `completed_at` for tasks

If easy, we can later show simple stats in an internal admin view or just query the DB.

6. Non-goals (Out of Scope for v0.1)

To keep focus, we explicitly **exclude** for now:

- EMR/HIS integrations
- Automatic data ingestion (labs, vitals, notes)
- AI/LLM-based “what changed?”
- State-aware modes (pre-round / bedside / handoff)
- Notifications / reminders (push / SMS / email) – unless you think a very simple version is trivial
- Team collaboration features between multiple doctors

v0.1 = **single doctor, single device, manual data entry**, but with a **much smarter queue** than their current mental one.

7. Tech & Implementation – High-level (up to you)

This part is mostly your call as tech lead, but my preferences:

- Web app, mobile-first design

- Any modern stack you're comfortable moving fast with (React / Next.js / etc. for frontend, Node / Python / etc. for backend)
- Simple deployment (e.g. Vercel / Render / Railway)
- Shared task board in Notion / Trello for sprint tracking

I don't care about stack purity in v0.1 – I care about **speed, stability, and ability to iterate.**

8. 4-week Sprint Breakdown (Proposal)

You can adjust based on your sense of effort, but here's a first pass:

Week 1 – Skeleton & Data Model

- Set up repo + basic project structure
- Implement:
 - Data model for User / Patient / Task
 - Very simple storage (DB or even hosted DB like Supabase)
- Rough UI for:
 - Patients list (no fancy design)
 - Patient detail
- No full auth yet or very basic auth
- Internal-only version for us to click through

Output:

Clickable prototype with static-ish data so we can validate flow.

Week 2 – Today Queue + CRUD

- Implement **Today Queue** screen with sorting logic
- Full CRUD:
 - Create/Edit/Delete patient
 - Create/Edit/Complete task
- Connect queue ↔ patient detail
- Make UI at least **clean and usable on mobile** (no need for perfect design)

Output:

First end-to-end version where a test doctor could:

- Add patients
- Add tasks
- See them in the queue
- Use it for a few minutes continuously

Week 3 – Polishing & Test Users

- Basic polishing:
 - Better layout on mobile
 - Small improvements in UX (shortcuts, default values)
- Add manual “What changed” notes on patient detail
- Add minimal tracking (last_login, tasks created/completed)

On my side (Nour):

- Bring in **2–3 real doctors** to test the app in real life
 - Collect structured feedback (what’s confusing, what’s missing, what’s useful)
-

Week 4 – Iterate on Real Feedback

Based on feedback from first users:

- Fix obvious bugs / friction points
- Add **one or two** high-impact improvements, for example:
 - Quick “add task” shortcut from queue
 - Simple filter (e.g. “show only critical patients”)
 - Better default due times (“today evening”, “tomorrow morning”)

By end of week 4:

- We should have:
 - At least 3–5 doctors who have used it multiple times
 - Concrete quotes + metrics
 - A clear sense if this is worth going deeper
-

9. Success Criteria for the 4-week Test

We consider the 4-week sprint successful if:

- Product:
 - v0.1 is live, stable enough to use in real shifts
- Users:
 - 3–5 doctors have used it for at least **3 days** each
 - At least 2 of them say they’d be unhappy if we took it away
- Team:
 - We both feel that:
 - We can work well together
 - We move at a reasonable speed given our constraints
 - It’s worth discussing deeper commitment / equity / next steps

If not, we still win because:

- We’ll have tested the idea in reality

- We'll have code, structure, and learnings for the next iteration
 - We'll know if we should pivot, change team, or change approach
-